

Assignment 4: AutoML - Hyperparameter Optimization

This notebook implements two hyperparameter optimization strategies for a neural network classifier:

1. **Genetic Algorithm** - Evolutionary approach with roulette selection, one-point crossover, and age-based replacement
2. **Bayesian Optimization** - Probabilistic model-based optimization

Hyperparameters being tuned:

- Mini-batch size $B \in [16, 1024]$
- Activation function $\in \{\text{ReLU}, \text{Sigmoid}, \text{Tanh}\}$

```
# Core dependencies
import numpy as np
import torch
import torch.nn as nn
from torch.utils.data import DataLoader, TensorDataset, Subset
from sklearn.metrics import f1_score
import matplotlib.pyplot as plt
from bayes_opt import BayesianOptimization
import random
import copy

# Reproducibility settings
RANDOM_SEED = 42
random.seed(RANDOM_SEED)
np.random.seed(RANDOM_SEED)
torch.manual_seed(RANDOM_SEED)

# Compute device selection
COMPUTE_DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {COMPUTE_DEVICE}")

Using device: cuda
```

Data Preparation

Load the Federated EMNIST dataset, filter for digit classes (0-9), normalize, and create train/validation/test splits.

```
def prepare_dataset():
    """
```

```

Loads and preprocesses the EMNIST data:
- Extracts images/labels from federated format
- Filters to keep only digit classes
- Normalizes pixel values
- Creates 80/20 train/validation split
"""

def unpack_federated_data(file_path):
    """Extract all client data from federated format into single arrays."""
    raw = np.load(file_path, allow_pickle=True)
    if raw.ndim == 0:
        raw = raw.item()

    images_list, labels_list = [], []
    data_iterator = raw.values() if isinstance(raw, dict) else raw

    for client_data in data_iterator:
        if isinstance(client_data, dict) and 'images' in client_data:
            images_list.append(client_data['images'])
            labels_list.append(client_data['labels'])

    return np.vstack(images_list), np.hstack(labels_list)

# Load raw federated data
X_train_raw, y_train_raw =
unpack_federated_data('data/train_data.npy')
X_test_raw, y_test_raw =
unpack_federated_data('data/test_data.npy')

print(f"Raw training samples: {len(y_train_raw)}")

# Keep only digit classes (labels 0-9)
digit_mask_train = y_train_raw < 10
digit_mask_test = y_test_raw < 10

X_train_digits = X_train_raw[digit_mask_train]
y_train_digits = y_train_raw[digit_mask_train]
X_test_digits = X_test_raw[digit_mask_test]
y_test_digits = y_test_raw[digit_mask_test]

print(f"Digit-only training samples: {len(y_train_digits)}")

# Remap labels to contiguous range [0, num_classes-1]
unique_classes = np.unique(y_train_digits)
label_mapping = {old_label: new_idx for new_idx, old_label in
enumerate(sorted(unique_classes))}

y_train_mapped = np.array([label_mapping[	lbl] for lbl in

```

```

y_train_digits])
    y_test_mapped = np.array([label_mapping[lbl] for lbl in
y_test_digits])

# Normalize pixel intensities to [0, 1]
X_train_norm = X_train_digits.astype(np.float32)
X_test_norm = X_test_digits.astype(np.float32)
if X_train_norm.max() > 1.0:
    X_train_norm /= 255.0
    X_test_norm /= 255.0

# Flatten images if multi-dimensional
if X_train_norm.ndim > 2:
    X_train_norm = X_train_norm.reshape(X_train_norm.shape[0], -1)
    X_test_norm = X_test_norm.reshape(X_test_norm.shape[0], -1)

# Convert to PyTorch tensors
X_train_tensor = torch.tensor(X_train_norm)
y_train_tensor = torch.tensor(y_train_mapped, dtype=torch.long)
X_test_tensor = torch.tensor(X_test_norm)
y_test_tensor = torch.tensor(y_test_mapped, dtype=torch.long)

# Create full training dataset
full_train_dataset = TensorDataset(X_train_tensor, y_train_tensor)
test_dataset = TensorDataset(X_test_tensor, y_test_tensor)

# 80/20 stratified-ish split using shuffled indices
n_total = len(full_train_dataset)
n_train = int(0.8 * n_total)

shuffled_indices = np.random.permutation(n_total)
train_indices = shuffled_indices[:n_train]
val_indices = shuffled_indices[n_train:]

train_subset = Subset(full_train_dataset, train_indices)
val_subset = Subset(full_train_dataset, val_indices)

feature_dim = X_train_norm.shape[1]
n_classes = len(unique_classes)

print(f"Feature dimension: {feature_dim}")
print(f"Number of classes: {n_classes}")
print(f"Training samples: {len(train_subset)}, Validation samples: {len(val_subset)}, Test samples: {len(test_dataset)})")

return train_subset, val_subset, test_dataset, feature_dim,
n_classes

# Execute data preparation

```

```

train_data, val_data, test_data, INPUT_DIM, NUM_CLASSES =
prepare_dataset()

Raw training samples: 31825
Digit-only training samples: 9517
Feature dimension: 784
Number of classes: 10
Training samples: 7613, Validation samples: 1904, Test samples: 1049

```

Neural Network Architecture

A two-layer feedforward network with a configurable hidden layer activation function.

```

class FeedForwardClassifier(nn.Module):
    """
        Simple two-layer neural network for classification.
        Architecture: Input -> Hidden (with activation) -> Output
    """
    ACTIVATION_MAP = {
        'relu': nn.ReLU,
        'sigmoid': nn.Sigmoid,
        'tanh': nn.Tanh
    }

    def __init__(self, in_features, hidden_units, out_features,
activation_type='relu'):
        super().__init__()
        self.hidden_layer = nn.Linear(in_features, hidden_units)
        self.output_layer = nn.Linear(hidden_units, out_features)

        # Get activation function from mapping, default to ReLU
        activation_class =
self.ACTIVATION_MAP.get(activation_type.lower(), nn.ReLU)
        self.activation_fn = activation_class()

    def forward(self, x):
        hidden_out = self.activation_fn(self.hidden_layer(x))
        return self.output_layer(hidden_out)

def compute_validation_f1(batch_sz, activation_type, training_set,
validation_set,
                           in_dim, n_classes, num_epochs=10):
    """
        Train a model with given hyperparameters and return validation F1
score.
        Uses SGD optimizer with mini-batch training.
    """
    # Clamp batch size to valid range

```

```

batch_sz = int(np.clip(batch_sz, 16, 1024))

train_loader = DataLoader(training_set, batch_size=batch_sz,
shuffle=True)
val_loader = DataLoader(validation_set, batch_size=batch_sz,
shuffle=False)

# Initialize model, optimizer, and loss function
model = FeedForwardClassifier(in_dim, 256, n_classes,
activation_type).to(COMPUTE_DEVICE)
sgd_optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
loss_fn = nn.CrossEntropyLoss()

# Training loop
for _ in range(num_epochs):
    model.train()
    for batch_x, batch_y in train_loader:
        batch_x = batch_x.to(COMPUTE_DEVICE)
        batch_y = batch_y.to(COMPUTE_DEVICE)

        sgd_optimizer.zero_grad()
        predictions = model(batch_x)
        loss = loss_fn(predictions, batch_y)
        loss.backward()
        sgd_optimizer.step()

# Validation evaluation
model.eval()
all_predictions = []
all_ground_truth = []

with torch.no_grad():
    for batch_x, batch_y in val_loader:
        batch_x = batch_x.to(COMPUTE_DEVICE)
        logits = model(batch_x)
        predicted_classes = torch.argmax(logits, dim=1)
        all_predictions.extend(predicted_classes.cpu().numpy())
        all_ground_truth.extend(batch_y.numpy())

    return f1_score(all_ground_truth, all_predictions,
average='macro')

```

Part 1: Genetic Algorithm

Implementation featuring:

- **Roulette wheel selection:** Parents chosen proportionally to their fitness
- **One-point crossover:** Genetic material split at a random point between parents

- **Age-based replacement:** Individuals removed after surviving a maximum number of generations
- **Mutation:** Random perturbations to maintain diversity

```

class EvolutionaryHyperparamSearch:
    """
        Genetic Algorithm for hyperparameter optimization with:
        - Roulette wheel parent selection
        - One-point crossover
        - Age-based population management
    """

    AVAILABLE_ACTIVATIONS = ['relu', 'sigmoid', 'tanh']

    def __init__(self, population_count, generation_limit,
max_lifespan,
                 train_set, val_set, input_dim, num_classes,
mutation_rate=0.15):
        self.population_count = population_count
        self.generation_limit = generation_limit
        self.max_lifespan = max_lifespan # Maximum generations an
individual can survive
        self.mutation_rate = mutation_rate

        self.train_set = train_set
        self.val_set = val_set
        self.input_dim = input_dim
        self.num_classes = num_classes

        self.current_population = []
        self.generation_best_scores = []
        self.generation_avg_scores = []

    def _create_individual(self):
        """Generate a random individual with age 0."""
        return {
            'batch_size': random.randint(16, 1024),
            'activation': random.choice(self.AVAILABLE_ACTIVATIONS),
            'fitness': 0.0,
            'age': 0
        }

    def _evaluate_fitness(self, individual):
        """Compute fitness (F1 score) for an individual."""
        score = compute_validation_f1(
            individual['batch_size'],
            individual['activation'],
            self.train_set,
            self.val_set,
            self.input_dim,

```

```

        self.num_classes
    )
individual['fitness'] = score
return score

def _roulette_wheel_select(self):
    """
    Select a parent using roulette wheel (fitness-proportionate)
    selection.
    Higher fitness = higher probability of selection.
    """
    fitness_values = [ind['fitness'] for ind in
self.current_population]
    total_fitness = sum(fitness_values)

    # Handle edge case where all fitness is zero
    if total_fitness <= 0:
        return random.choice(self.current_population)

    # Spin the wheel
    spin = random.uniform(0, total_fitness)
    cumulative = 0

    for individual in self.current_population:
        cumulative += individual['fitness']
        if cumulative >= spin:
            return individual

    return self.current_population[-1]

def _one_point_crossover(self, parent_a, parent_b):
    """
    Perform one-point crossover between two parents.
    Genes: [batch_size, activation]
    Crossover point determines which genes come from which parent.
    """
    # For 2 genes, crossover point can be 0 or 1
    # Point 0: child gets all from parent_a
    # Point 1: child gets first gene from parent_a, second from
parent_b
    crossover_point = random.randint(0, 1)

    if crossover_point == 0:
        # All genes from parent_a
        child_batch = parent_a['batch_size']
        child_activation = parent_a['activation']
    else:
        # First gene from parent_a, second from parent_b
        child_batch = parent_a['batch_size']

```

```

        child_activation = parent_b['activation']

    return {
        'batch_size': child_batch,
        'activation': child_activation,
        'fitness': 0.0,
        'age': 0 # New offspring starts at age 0
    }

def _apply_mutation(self, individual):
    """Apply random mutations based on mutation rate."""
    if random.random() < self.mutation_rate:
        individual['batch_size'] = random.randint(16, 1024)

    if random.random() < self.mutation_rate:
        individual['activation'] =
random.choice(self.AVAILABLE_ACTIVATIONS)

    return individual

def _age_based_culling(self):
    """Remove individuals that have exceeded maximum lifespan."""
    self.current_population = [
        ind for ind in self.current_population
        if ind['age'] < self.max_lifespan
    ]

def evolve(self):
    """Run the evolutionary process."""
    # Initialize population
    self.current_population = [self._create_individual() for _ in
range(self.population_count)]

    # Evaluate initial population
    print("Evaluating initial population...")
    for ind in self.current_population:
        self._evaluate_fitness(ind)

    # Evolution loop
    for gen in range(self.generation_limit):
        # Age all individuals
        for ind in self.current_population:
            ind['age'] += 1

        # Remove individuals that are too old
        self._age_based_culling()

        # Sort by fitness (descending)
        self.current_population.sort(key=lambda x: x['fitness'],

```

```

reverse=True)

        # Record statistics
        current_best = self.current_population[0]['fitness']
        current_avg = np.mean([ind['fitness'] for ind in
self.current_population])
            self.generation_best_scores.append(current_best)
            self.generation_avg_scores.append(current_avg)

        print(f"Generation {gen + 1}/{self.generation_limit} | "
              f"Best: {current_best:.4f} | Avg: {current_avg:.4f}
| "
              f"Pop size: {len(self.current_population)})"

    # Build next generation
    next_generation = []

    # Elitism: keep the best individual (with age preserved)
    elite = copy.deepcopy(self.current_population[0])
    next_generation.append(elite)

    # Generate offspring to fill population
    while len(next_generation) < self.population_count:
        # Select parents via roulette wheel
        mom = self._roulette_wheel_select()
        dad = self._roulette_wheel_select()

        # Create offspring via one-point crossover
        child = self._one_point_crossover(mom, dad)

        # Apply mutation
        child = self._apply_mutation(child)

        # Evaluate offspring fitness
        self._evaluate_fitness(child)

        next_generation.append(child)

    self.current_population = next_generation

    # Return best individual from final population
    self.current_population.sort(key=lambda x: x['fitness'],
reverse=True)
    return self.current_population[0]

def plot_evolution_history(self):
    """Visualize fitness progression across generations."""
    generations = range(1, len(self.generation_best_scores) + 1)

```

```

        plt.figure(figsize=(10, 6))
        plt.plot(generations, self.generation_avg_scores, 'b-o',
label='Average Fitness', alpha=0.7)
        plt.plot(generations, self.generation_best_scores, 'r-s',
label='Best Fitness', alpha=0.7)
        plt.xlabel('Generation')
        plt.ylabel('Validation F1 Score (Macro)')
        plt.title('Genetic Algorithm: Fitness Evolution')
        plt.legend()
        plt.grid(True, alpha=0.3)
        plt.tight_layout()
        plt.show()

# Run Genetic Algorithm search
ga_searcher = EvolutionaryHyperparamSearch(
    population_count=20,
    generation_limit=20,
    max_lifespan=5, # Individuals removed after 5 generations
    train_set=train_data,
    val_set=val_data,
    input_dim=INPUT_DIM,
    num_classes=NUM_CLASSES,
    mutation_rate=0.15
)

best_ga_solution = ga_searcher.evolve()

print("\n" + "*50")
print("GENETIC ALGORITHM RESULTS")
print("*50")
print(f"Best Batch Size: {best_ga_solution['batch_size']}")"
print(f"Best Activation: {best_ga_solution['activation']}")"
print(f"Validation F1 Score: {best_ga_solution['fitness']:.4f}")"

Evaluating initial population...
Generation 1/20 | Best: 0.7755 | Avg: 0.3038 | Pop size: 20
Generation 2/20 | Best: 0.7905 | Avg: 0.4863 | Pop size: 20
Generation 3/20 | Best: 0.7905 | Avg: 0.5764 | Pop size: 20
Generation 4/20 | Best: 0.7946 | Avg: 0.5829 | Pop size: 20
Generation 5/20 | Best: 0.7946 | Avg: 0.6516 | Pop size: 20
Generation 6/20 | Best: 0.8101 | Avg: 0.6409 | Pop size: 20
Generation 7/20 | Best: 0.8101 | Avg: 0.6471 | Pop size: 20
Generation 8/20 | Best: 0.8101 | Avg: 0.5678 | Pop size: 20
Generation 9/20 | Best: 0.8101 | Avg: 0.7075 | Pop size: 20
Generation 10/20 | Best: 0.7871 | Avg: 0.6713 | Pop size: 19
Generation 11/20 | Best: 0.7965 | Avg: 0.6827 | Pop size: 20
Generation 12/20 | Best: 0.7965 | Avg: 0.6112 | Pop size: 20
Generation 13/20 | Best: 0.7965 | Avg: 0.5711 | Pop size: 20
Generation 14/20 | Best: 0.7994 | Avg: 0.7392 | Pop size: 20
Generation 15/20 | Best: 0.7994 | Avg: 0.6603 | Pop size: 20

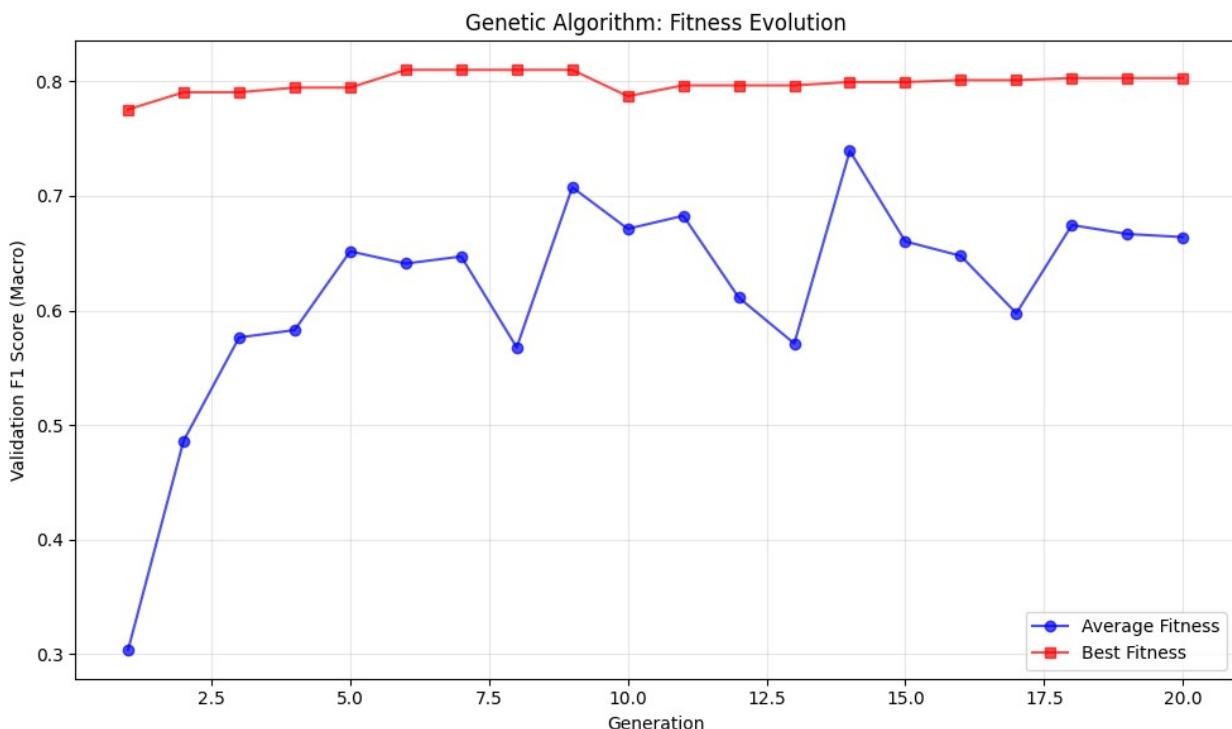
```

Generation 16/20	Best: 0.8011	Avg: 0.6476	Pop size: 20
Generation 17/20	Best: 0.8011	Avg: 0.5975	Pop size: 20
Generation 18/20	Best: 0.8029	Avg: 0.6745	Pop size: 20
Generation 19/20	Best: 0.8029	Avg: 0.6667	Pop size: 20
Generation 20/20	Best: 0.8029	Avg: 0.6641	Pop size: 20

GENETIC ALGORITHM RESULTS

Best Batch Size: 41
 Best Activation: tanh
 Validation F1 Score: 0.8029

```
# Plot GA evolution progress
ga_searcher.plot_evolution_history()
```



Part 2: Bayesian Optimization

Using the `bayesian-optimization` package to find optimal hyperparameters by modeling the objective function with a Gaussian Process surrogate.

```
def bayesian_hyperparameter_search(train_set, val_set, in_dim,
n_classes,
n_iterations=40):
    ...
```

Use Bayesian Optimization to find optimal batch size and activation function.

Activation is encoded as a continuous value [0, 2] and discretized.

```

"""
activation_options = ['relu', 'sigmoid', 'tanh']

def objective_function(batch_size_continuous,
activation_continuous):
    """Black-box function: returns validation F1 for given
hyperparameters."""
    # Discretize activation index
    act_idx = int(np.clip(np.round(activation_continuous), 0, 2))
    chosen_activation = activation_options[act_idx]

    return compute_validation_f1(
        batch_size_continuous,
        chosen_activation,
        train_set,
        val_set,
        in_dim,
        n_classes
    )

# Define search space bounds
search_bounds = {
    'batch_size_continuous': (16, 1024),
    'activation_continuous': (0, 2)
}

# Initialize optimizer
bo_optimizer = BayesianOptimization(
    f=objective_function,
    pbounds=search_bounds,
    random_state=RANDOM_SEED,
    verbose=2
)

# Run optimization
bo_optimizer.maximize(init_points=n_random_init,
n_iter=n_iterations)

# Extract best parameters
optimal_params = bo_optimizer.max['params']
optimal_batch = int(optimal_params['batch_size_continuous'])
optimal_act_idx =
int(np.round(optimal_params['activation_continuous']))
optimal_activation = activation_options[optimal_act_idx]

```

```

        return optimal_batch, optimal_activation,
bo_optimizer.max['target']

# Execute Bayesian Optimization
print("Starting Bayesian Optimization...")
bo_batch_size, bo_activation, bo_best_f1 =
bayesian_hyperparameter_search(
    train_data, val_data, INPUT_DIM, NUM_CLASSES
)

print("\n" + "*50)
print("BAYESIAN OPTIMIZATION RESULTS")
print("*50)
print(f"Best Batch Size: {bo_batch_size}")
print(f"Best Activation: {bo_activation}")
print(f"Validation F1 Score: {bo_best_f1:.4f}")

```

Starting Bayesian Optimization...

iter	target	batch_...	activa...
1	0.3534117	393.53643	1.9014286
2	0.0212206	753.84989	1.1973169
3	0.0205466	173.26678	0.3119890
4	0.7420676	74.548281	1.7323522
5	0.0212206	621.92393	1.4161451
6	0.7909241	36.749170	1.9398197
7	0.1480694	855.10218	0.4246782
8	0.5015409	199.27956	0.3668090
9	0.0212406	322.67618	1.0495128
10	0.0212206	451.40057	0.5824582
11	0.7510345	54.964284	0.0
12	0.3231825	37.629367	1.4176168
13	0.7466610	36.404013	1.7837637
14	0.3077065	55.006723	1.0848428
15	0.1559794	74.464319	0.8440364
16	0.7340078	75.045509	2.0
17	0.7671044	54.166285	0.0
18	0.7591630	53.084508	0.0
19	0.1239734	52.954762	1.0648687
20	0.6364867	73.687397	2.0
21	0.7444461	56.162823	0.0
22	0.7291477	57.361193	0.0
23	0.3775143	57.169302	1.1610528
24	0.6944234	58.652259	0.0
25	0.7405458	60.082062	0.0
26	0.1235297	59.711829	1.1896451
27	0.7211617	61.179419	0.0
28	0.7475639	62.520715	0.0
29	0.1589152	62.304431	1.1457443
30	0.7643962	63.596931	0.0

40	0.7754090	33.888017	2.0
31	0.7057267	64.791051	0.0
32	0.0834684	64.595679	1.1921785
33	0.6986485	66.050479	0.0
34	0.7277695	67.310517	0.0
35	0.1273620	67.111469	1.1251365
36	0.7186683	68.353958	0.0
37	0.5422197	69.534213	0.0
38	0.6784292	35.128979	2.0
39	0.3785942	35.189947	0.7699149
41	0.7893263	32.855774	2.0
42	0.3439113	33.111855	0.9462515
43	0.7197047	31.854245	2.0
44	0.7172699	30.527363	2.0
45	0.3582978	30.823651	0.8567877
46	0.7921513	29.391197	2.0
47	0.7984008	28.266424	2.0
48	0.3881346	28.484727	0.8241096
49	0.7445813	27.157783	2.0
50	0.8209178	25.701563	2.0

BAYESIAN OPTIMIZATION RESULTS

Best Batch Size: 25
 Best Activation: tanh
 Validation F1 Score: 0.8209

Final Model Training and Evaluation

Train models using the best hyperparameters from each method on the combined train+validation data, then evaluate on the held-out test set.

```
def train_final_model(batch_sz, activation, train_set, val_set,
test_set,
                     in_dim, n_classes, total_epochs=50,
method_label=""):
    """
        Train a model on combined train+val data with selected
hyperparameters.
        Track training F1 per epoch and evaluate on test set.
    """
    print(f"\n{'='*60}")
    print(f"Final Training: {method_label}")
    print(f"Hyperparameters: batch_size={batch_sz},
activation={activation}")
    print('='*60)
```

```

# Combine training and validation data
combined_dataset = torch.utils.data.ConcatDataset([train_set,
val_set])
combined_loader = DataLoader(combined_dataset,
batch_size=batch_sz, shuffle=True)
test_loader = DataLoader(test_set, batch_size=batch_sz,
shuffle=False)

# Initialize fresh model
model = FeedForwardClassifier(in_dim, 256, n_classes,
activation).to(COMPUTE_DEVICE)
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
loss_function = nn.CrossEntropyLoss()

epoch_f1_scores = []

for epoch in range(total_epochs):
    model.train()
    epoch_predictions = []
    epoch_labels = []

    for batch_x, batch_y in combined_loader:
        batch_x = batch_x.to(COMPUTE_DEVICE)
        batch_y = batch_y.to(COMPUTE_DEVICE)

        optimizer.zero_grad()
        outputs = model(batch_x)
        loss = loss_function(outputs, batch_y)
        loss.backward()
        optimizer.step()

    # Collect predictions for F1 calculation
    preds = torch.argmax(outputs, dim=1)
    epoch_predictions.extend(preds.cpu().numpy())
    epoch_labels.extend(batch_y.cpu().numpy())

    # Calculate training F1 for this epoch
    train_f1 = f1_score(epoch_labels, epoch_predictions,
average='macro')
    epoch_f1_scores.append(train_f1)

    if (epoch + 1) % 10 == 0:
        print(f"Epoch {epoch + 1:3d}/{total_epochs}: Training F1 = {train_f1:.4f}")

    # Evaluate on test set
    model.eval()
    test_predictions = []
    test_labels = []

```

```

with torch.no_grad():
    for batch_x, batch_y in test_loader:
        batch_x = batch_x.to(COMPUTE_DEVICE)
        outputs = model(batch_x)
        preds = torch.argmax(outputs, dim=1)
        test_predictions.extend(preds.cpu().numpy())
        test_labels.extend(batch_y.numpy())

test_f1 = f1_score(test_labels, test_predictions, average='macro')
print(f"\nFinal Test F1 Score: {test_f1:.4f}")

return epoch_f1_scores, test_f1

# Train final model using GA-selected hyperparameters
ga_training_history, ga_test_score = train_final_model(
    best_ga_solution['batch_size'],
    best_ga_solution['activation'],
    train_data, val_data, test_data,
    INPUT_DIM, NUM_CLASSES,
    total_epochs=50,
    method_label="Genetic Algorithm"
)

```

```
=====
Final Training: Genetic Algorithm
Hyperparameters: batch_size=41, activation=tanh
=====
```

```

Epoch 10/50: Training F1 = 0.8095
Epoch 20/50: Training F1 = 0.8500
Epoch 30/50: Training F1 = 0.8702
Epoch 40/50: Training F1 = 0.8813
Epoch 50/50: Training F1 = 0.8846

```

```
Final Test F1 Score: 0.8826
```

```

# Train final model using BO-selected hyperparameters
bo_training_history, bo_test_score = train_final_model(
    bo_batch_size,
    bo_activation,
    train_data, val_data, test_data,
    INPUT_DIM, NUM_CLASSES,
    total_epochs=50,
    method_label="Bayesian Optimization"
)

```

```
=====
Final Training: Bayesian Optimization
=====
```

```

Hyperparameters: batch_size=25, activation=tanh
=====
Epoch 10/50: Training F1 = 0.8255
Epoch 20/50: Training F1 = 0.8625
Epoch 30/50: Training F1 = 0.8798
Epoch 40/50: Training F1 = 0.8876
Epoch 50/50: Training F1 = 0.8940

Final Test F1 Score: 0.9008

# Plot training curves for both methods
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

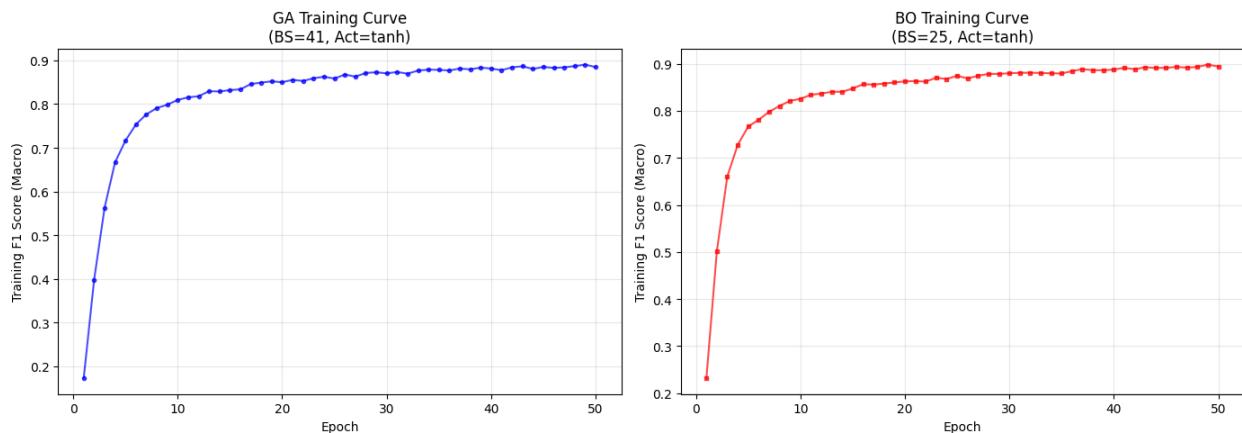
epochs_range = range(1, 51)

# GA Training Curve
axes[0].plot(epochs_range, ga_training_history, 'b-o', markersize=3,
alpha=0.7)
axes[0].set_xlabel('Epoch')
axes[0].set_ylabel('Training F1 Score (Macro)')
axes[0].set_title(f'GA Training Curve\nn(BS={best_ga_solution["batch_size"]}],\nAct={best_ga_solution["activation"]})')
axes[0].grid(True, alpha=0.3)

# BO Training Curve
axes[1].plot(epochs_range, bo_training_history, 'r-s', markersize=3,
alpha=0.7)
axes[1].set_xlabel('Epoch')
axes[1].set_ylabel('Training F1 Score (Macro)')
axes[1].set_title(f'BO Training Curve\nn(BS={bo_batch_size}],\nAct={bo_activation})')
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```



Results Summary and Comparison

```
# Final results comparison
print("*70)
print("FINAL RESULTS COMPARISON")
print("*70)
print(f"\n{'Method':<25} {'Batch Size':<15} {'Activation':<12} {'Test F1':<10}")
print("-*70)
print(f"{'Genetic Algorithm':<25} {best_ga_solution['batch_size']:<15}
{best_ga_solution['activation']:<12} {ga_test_score:.4f}")
print(f"{'Bayesian Optimization':<25} {bo_batch_size:<15}
{bo_activation:<12} {bo_test_score:.4f}")
print("*70)

=====
FINAL RESULTS COMPARISON
=====

Method           Batch Size      Activation    Test F1
-----
Genetic Algorithm     41            tanh        0.8826
Bayesian Optimization 25            tanh        0.9008
=====
```

Discussion: Genetic Algorithm vs Bayesian Optimization

Genetic Algorithm

Advantages:

- Explores diverse regions of the search space through population-based search
- Can escape local optima through mutation and crossover operators
- Naturally parallelizable (fitness evaluations can run concurrently)
- Does not require gradient information or assumptions about the objective function shape

Disadvantages:

- Computationally expensive due to large number of fitness evaluations
- Hyperparameter sensitive (population size, mutation rate, crossover type)
- May converge slowly compared to model-based approaches
- No principled way to balance exploration vs exploitation

Bayesian Optimization

Advantages:

- Sample-efficient: builds a probabilistic model to guide the search
- Principled exploration-exploitation tradeoff via acquisition function

- Works well with small evaluation budgets
- Provides uncertainty estimates about the objective function

Disadvantages:

- Struggles with high-dimensional search spaces
- Overhead of fitting the Gaussian Process surrogate model
- Assumes smoothness of the objective function
- Handling categorical variables (like activation function) requires encoding tricks

Comparison

Both methods are effective for hyperparameter optimization, but they excel in different scenarios. Genetic Algorithms are better suited for highly non-convex or discontinuous search spaces with many local optima, while Bayesian Optimization is preferred when evaluations are expensive and the search space is relatively smooth and low-dimensional.